

«Batch merge and merge path sort»

*Calcul haute performance:
programmation et algorithmique avancées*

MAIN 5:
Hala Ben Ali
Paolo Conti

I – MERGE PATH AND SORT


Reference:

O. Green, R. McColl and D.A. Bader «**GPU Merge Path - A GPU Merging Algorithm**».
26th ACM International Conference on Supercomputing (ICS), San Servolo Island, Venice, Italy, June 25-29,
2012

summary:

- Sorting an array on GPU
- Batch merging various small arrays
- Proposing an application

Only one block:




STEPS:

- PARTITIONING
- FINDING THE MERGE PATH
- MERGING

PARTITIONING


PARTITIONING

Divide the work among the
blocks




PARTITIONING

Binary cross diagonal search for
Merge Path on the partition
points on
global memory



FIND MERGE PATH

In parallel and relying on *shared
memory*



Example :
3 blocks
3 threads

- **X_gpu**


Qx block0			Qx block 1			Qx block 2		
--------------	--	--	---------------	--	--	---------------	--	--

- **Y_gpu**

Qy block0			Qy block 1			Qy block 2		
--------------	--	--	---------------	--	--	---------------	--	--


FIND MERGE PATH

In parallel and relying on *shared
memory*




FIND MERGE PATH

Write the path on the *global memory*




FIND MERGE PATH

Cycles in parallel



MERGE

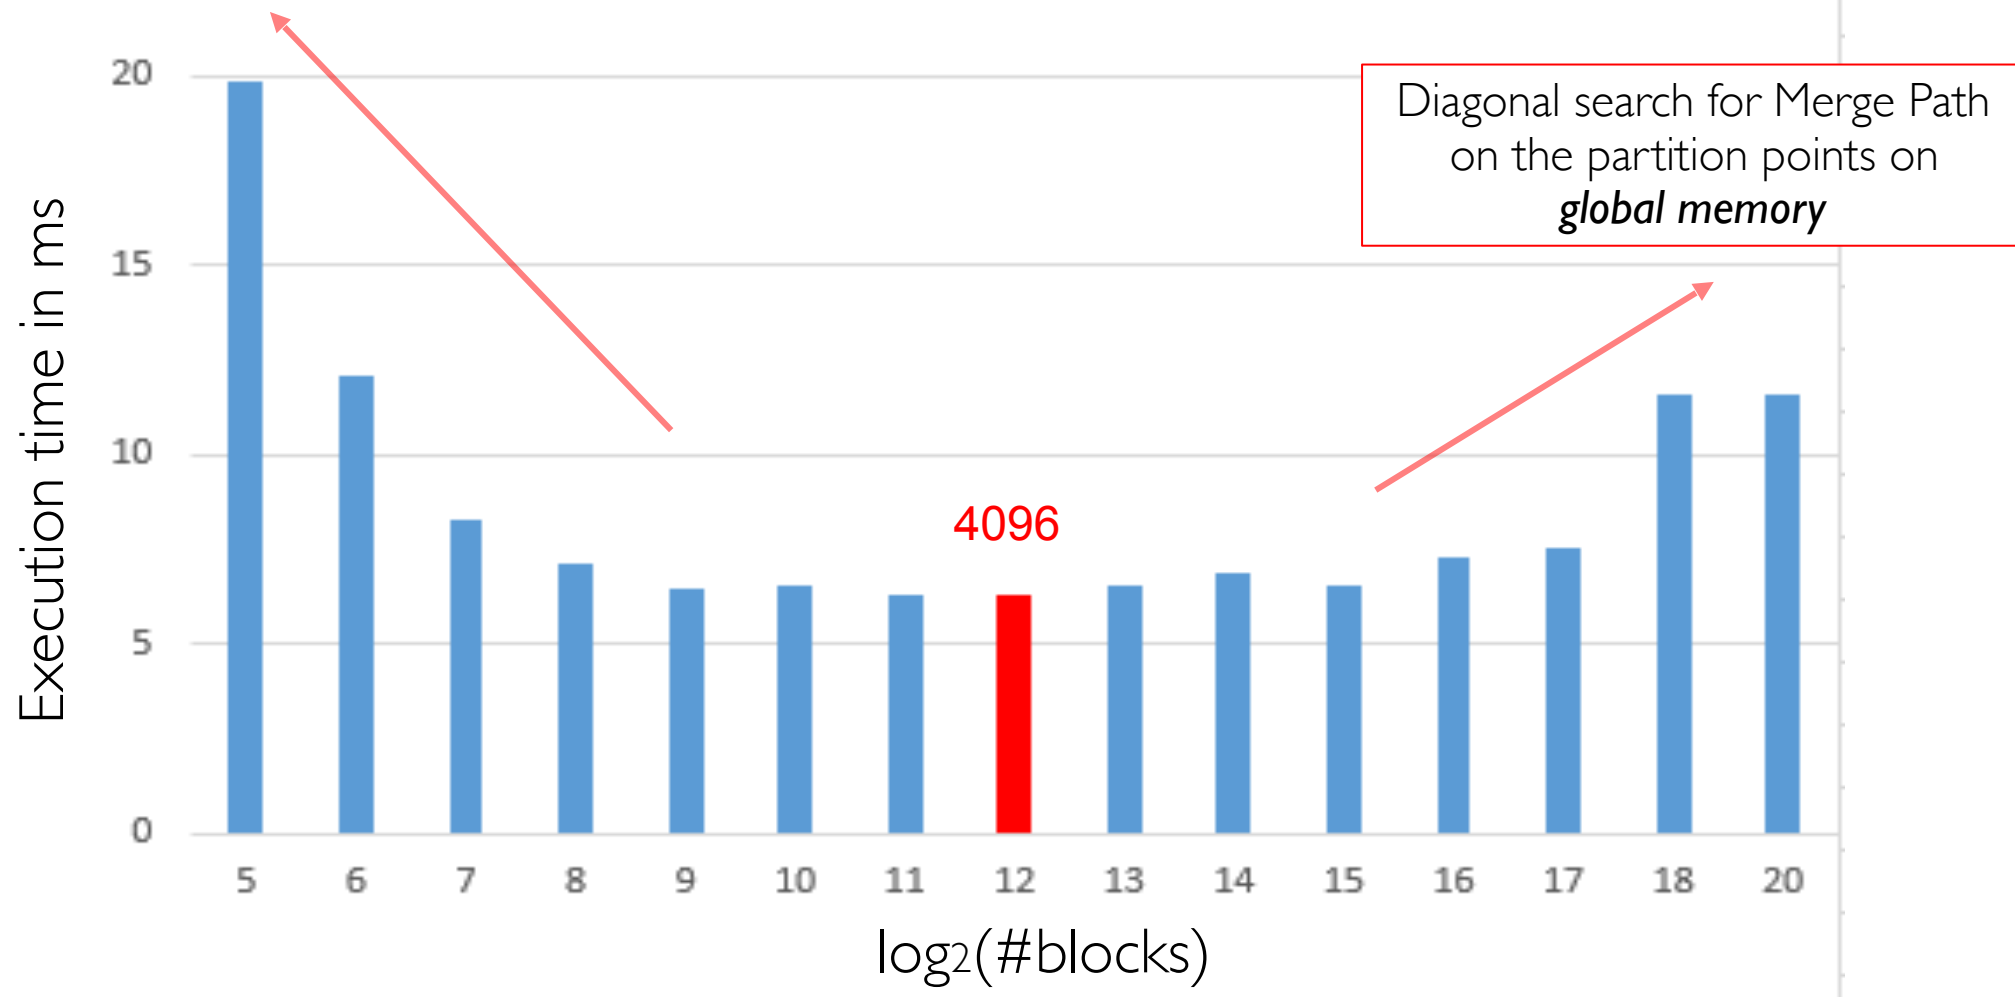
Placing the Merge Path on the *shared memory*



MERGE


Not enough blocks for a good parallelization

Scaling Number of Blocks with $N = 1,000,000$




EXECUTION TIMES

1-2



Our programme


1-3



Normal Merge Sort

2- BATCH MERGE

BATCH MERGE



INDEXES

d: number of elements in a pair of array


m: number of pairs of array mergesorted in a single block



```
int tidx = threadIdx.x % d;
```

```
int Qt = (threadIdx.x - tidx) / d;
```



```
int gbx = Qt + blockIdx.x * (blockDim.x / d);
```

- **Enumeration** from 0 to d-1 of the **elements** of **each of the m pairs** of arrays which are going to be sorted in a single block
- **Local enumeration** from 0 to m-1 of the pairs of arrays in a single block
- **Global enumeration** from 0 to N-1 of the pairs of arrays in all the blocks









MALLOC



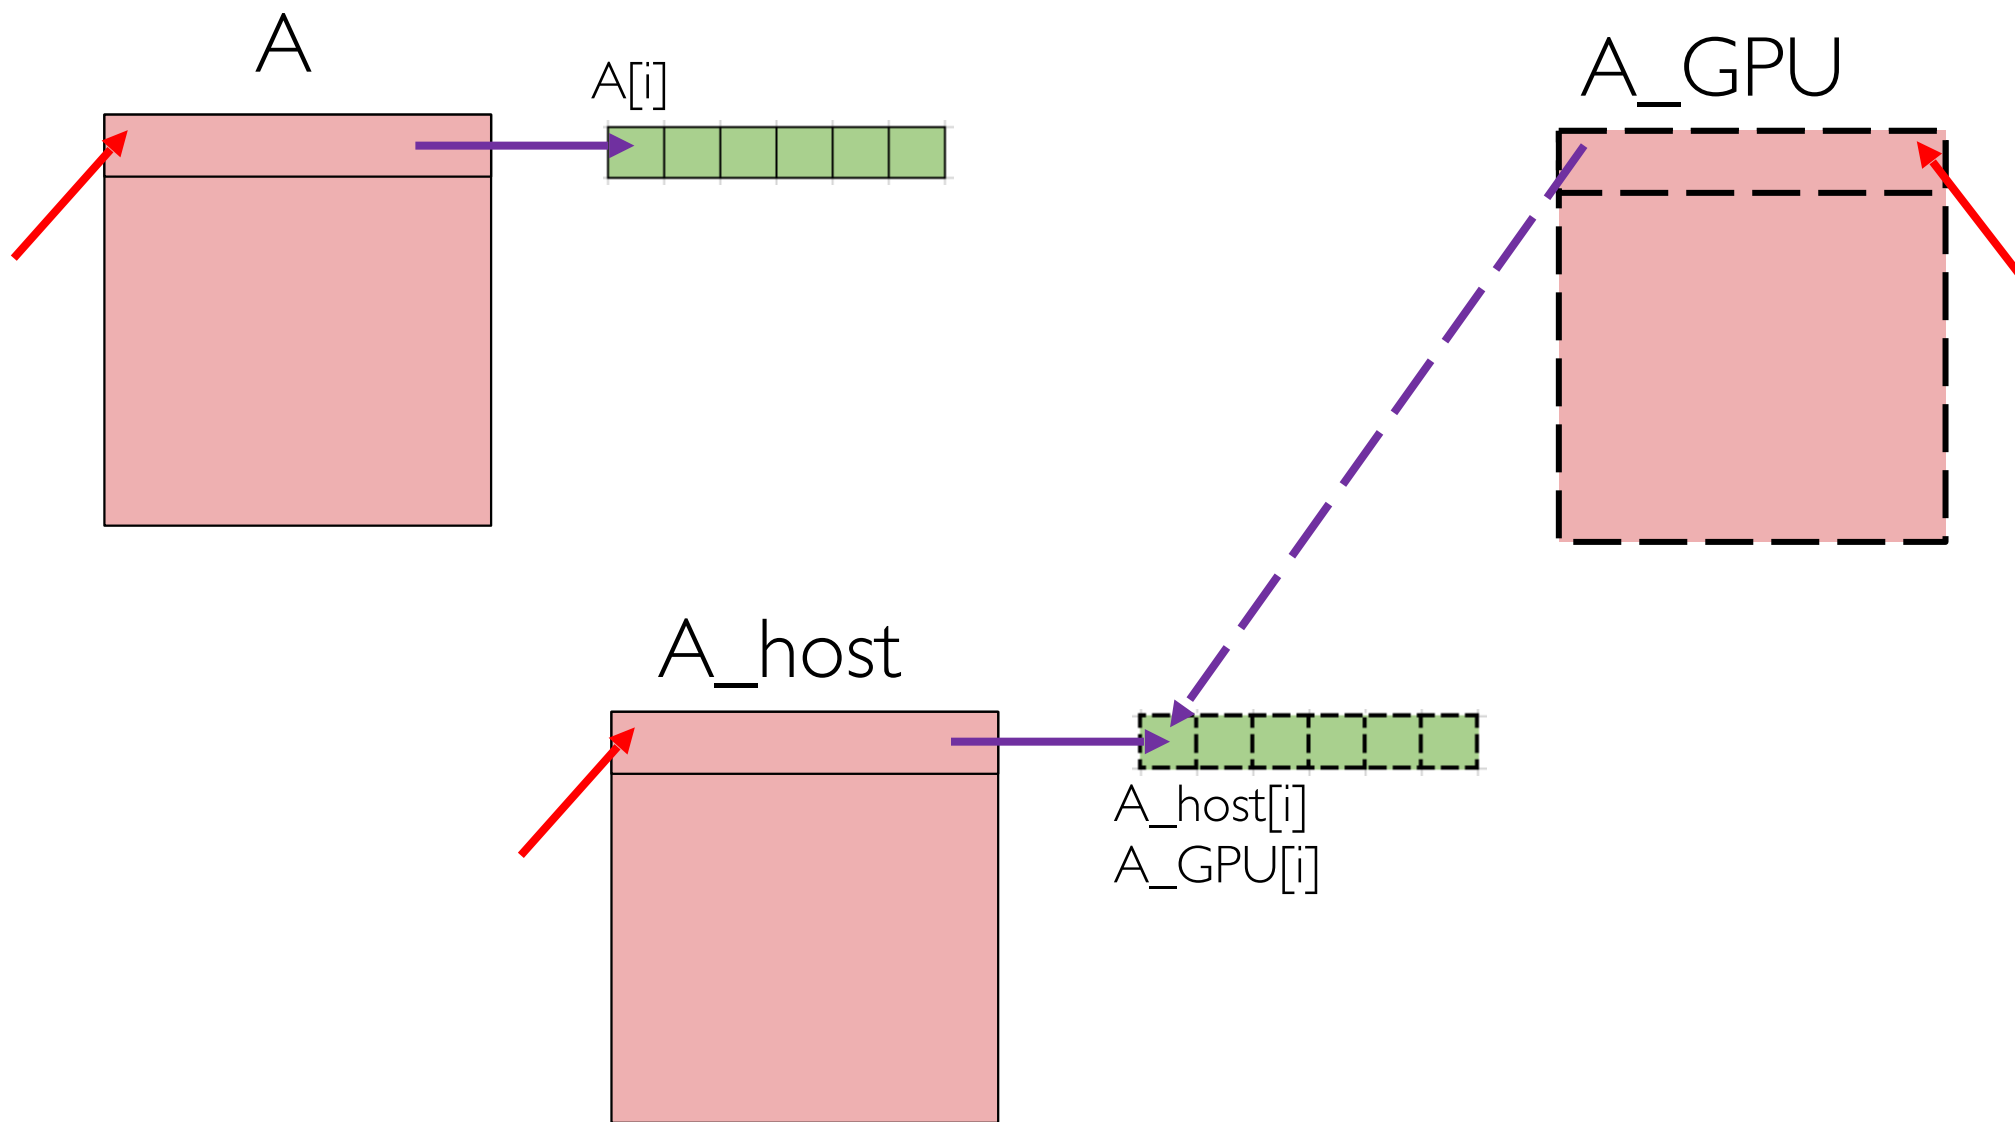
CUDA MALLOC



CUDA MEMCPY




CUDA MALLOC



CUDA MEMCPY


BATCH MERGE

$N = 1000$




BATCH MERGE

$N = 10,000$




BATCH MERGE

$N = 100.000$



BATCH MERGE

$N = 1.000.000$




N = 1000


VS

N = 1.000.000

N = 1e3



N = 1e6



3- BATCH MERGE APPLICATION


Merging «Spotify» Playlists



Merging
«*Spotify*»
playlists

GOAL

- GOAL: to obtain **new playlists** which merge users **favourite songs** with **new songs**



Merging
«*Spotify*»
playlists

DATA
STRUCT

- **SONG**

struct song{

- **Title**
- **Artist name**
- **Music genre**
- **Number of plays**

}

- **PLAYLIST**

Array of songs

Merging
«*Spotify*»
playlists

FEATURES
and
FUNCTIONS

- **UNIQUENESS:** No repetitions of the same song in the playlist.

Based on:

- Title
- Artist name

- **SORT-CRITERIUM:** Sorting according to a **coefficient** based on:

- **Music Genre**

Frequency of songs of the same musical genre in a playlist

- **Numer of plays**

-

Percentage with respect to the user total number of plays

- High number of plays
- Favourite genre

COEFFICIENT



- Low number of plays
- Less preferred genre

COEFFICIENT

